

```
In [1]: ### Project Topic  
  
# The purpose of this project is to classify breast cancer tumors as (benin or malignant).  
# This project will be using a simple multi-linear regression model.  
# Data source would come from the University of California, Irvine.  
# This project would help doctors and physicians who have trouble classifying mTBI especially at initial tests.
```

```
In [168]: ### Data  
  
# The data consist of an ID number, Diagnosis (M=malignant, B=Benign), and 30 features (tumor radius, texture, area, smoothness...)  
# 569 subjects would be used for this exercise.  
  
# Data was obtained from UCI's machine Learning repository  
# https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic  
  
import scipy as sp  
import scipy.stats as stats  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import copy  
# Set color map to have light blue background  
sns.set()  
import statsmodels.formula.api as smf  
import statsmodels.api as sm  
%matplotlib inline  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import roc_curve, auc, f1_score, roc_auc_score  
  
# Replace file_path with location of data.csv  
file_path = r'C:\Users\Eddie\Desktop\Engineering\wdbc_wheaders.csv'  
df = pd.read_csv(file_path)
```

```
In [169]: ### Data Cleaning - Display DF sample

# Data is initially .data file. File is processed externaly to add headers and
convert to a csv file.
# Realized data would be 570 x 32 matrix. A total of 30 features and 569 sampl
es.
# Data is clean and does not require much processing. None of the data has NAN
and data is label in an appropriate category.
# Only the diagnosis is cleaned to convert Benign to 0 and Malignant to 1.

#           ID      DIAGNOSIS    RADIUS1    TEXTURE1    PERIMETER1 ... FRACTAL_DI
MENSION3
# SUBJECT0
# SUBJECT1
# SUBJECT2
# SUBJECT3
# ...
# SUBJECT568

print(df.head()) # Display df
```

	ID	Diagnosis	radius1	texture1	perimeter1	area1	smoothness1	\
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	

	compactness1	concavity1	concave_points1	...	radius3	\
0	0.27760	0.3001	0.14710	...	25.38	
1	0.07864	0.0869	0.07017	...	24.99	
2	0.15990	0.1974	0.12790	...	23.57	
3	0.28390	0.2414	0.10520	...	14.91	
4	0.13280	0.1980	0.10430	...	22.54	

	texture3	perimeter3	area3	smoothness3	compactness3	concavity3	\
0	17.33	184.60	2019.0	0.1622	0.6656	0.7119	
1	23.41	158.80	1956.0	0.1238	0.1866	0.2416	
2	25.53	152.50	1709.0	0.1444	0.4245	0.4504	
3	26.50	98.87	567.7	0.2098	0.8663	0.6869	
4	16.67	152.20	1575.0	0.1374	0.2050	0.4000	

	concave_points3	symmetry3	fractal_dimension3
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678

[5 rows x 32 columns]

In [170]: *### Data Cleaning - Display DF features*

```
print(df.info()) # Display features
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
ID                    569 non-null int64
Diagnosis             569 non-null object
radius1              569 non-null float64
texture1             569 non-null float64
perimeter1           569 non-null float64
area1                569 non-null float64
smoothness1          569 non-null float64
compactness1         569 non-null float64
concavity1           569 non-null float64
concave_points1      569 non-null float64
symmetry1            569 non-null float64
fractal_dimension1   569 non-null float64
radius2              569 non-null float64
texture2             569 non-null float64
perimeter2           569 non-null float64
area2                569 non-null float64
smoothness2          569 non-null float64
compactness2         569 non-null float64
concavity2           569 non-null float64
concave_points2      569 non-null float64
symmetry2            569 non-null float64
fractal_dimension2   569 non-null float64
radius3              569 non-null float64
texture3             569 non-null float64
perimeter3           569 non-null float64
area3                569 non-null float64
smoothness3          569 non-null float64
compactness3         569 non-null float64
concavity3           569 non-null float64
concave_points3      569 non-null float64
symmetry3            569 non-null float64
fractal_dimension3   569 non-null float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.3+ KB
None
```

In [171]: *### Data Cleaning - Display Benign Count and Malignant Count*

```
diagnosis_counts = df['Diagnosis'].value_counts()
benign_count = diagnosis_counts['B']
malignant_count = diagnosis_counts['M']

print("Benign Count:", benign_count)
print("Malignant Count:", malignant_count)

df['Diagnosis'] = df['Diagnosis'].replace({'B': 0, 'M': 1})
```

```
Benign Count: 357
Malignant Count: 212
```

In [172]: *### Exploratory Data Analysis - display correlation matrix*


```
df.corr()
```

*## By observation of the correlation matrix it looks like radius, perimeter, and concave points are the most significant features.*

Out[172]:

	ID	Diagnosis	radius1	texture1	perimeter1	area1	smoothness
<b>ID</b>	1.000000	0.039769	0.074626	0.099770	0.073159	0.096893	-0.01296
<b>Diagnosis</b>	0.039769	1.000000	0.730029	0.415185	0.742636	0.708984	0.35856
<b>radius1</b>	0.074626	0.730029	1.000000	0.323782	0.997855	0.987357	0.17058
<b>texture1</b>	0.099770	0.415185	0.323782	1.000000	0.329533	0.321086	-0.02338
<b>perimeter1</b>	0.073159	0.742636	0.997855	0.329533	1.000000	0.986507	0.20727
<b>area1</b>	0.096893	0.708984	0.987357	0.321086	0.986507	1.000000	0.17702
<b>smoothness1</b>	-0.012968	0.358560	0.170581	-0.023389	0.207278	0.177028	1.00000
<b>compactness1</b>	0.000096	0.596534	0.506124	0.236702	0.556936	0.498502	0.65912
<b>concavity1</b>	0.050080	0.696360	0.676764	0.302418	0.716136	0.685983	0.52198
<b>concave_points1</b>	0.044158	0.776614	0.822529	0.293464	0.850977	0.823269	0.55369
<b>symmetry1</b>	-0.022114	0.330499	0.147741	0.071401	0.183027	0.151293	0.55777
<b>fractal_dimension1</b>	-0.052511	-0.012838	-0.311631	-0.076437	-0.261477	-0.283110	0.58479
<b>radius2</b>	0.143048	0.567134	0.679090	0.275869	0.691765	0.732562	0.30146
<b>texture2</b>	-0.007526	-0.008303	-0.097317	0.386358	-0.086761	-0.066280	0.06840
<b>perimeter2</b>	0.137331	0.556141	0.674172	0.281673	0.693135	0.726628	0.29609
<b>area2</b>	0.177742	0.548236	0.735864	0.259845	0.744983	0.800086	0.24655
<b>smoothness2</b>	0.096781	-0.067016	-0.222600	0.006614	-0.202694	-0.166777	0.33237
<b>compactness2</b>	0.033961	0.292999	0.206000	0.191975	0.250744	0.212583	0.31894
<b>concavity2</b>	0.055239	0.253730	0.194204	0.143293	0.228082	0.207660	0.24839
<b>concave_points2</b>	0.078768	0.408042	0.376169	0.163851	0.407217	0.372320	0.38067
<b>symmetry2</b>	-0.017306	-0.006522	-0.104321	0.009127	-0.081629	-0.072497	0.20077
<b>fractal_dimension2</b>	0.025725	0.077972	-0.042641	0.054458	-0.005523	-0.019887	0.28360
<b>radius3</b>	0.082405	0.776454	0.969539	0.352573	0.969476	0.962746	0.21312
<b>texture3</b>	0.064720	0.456903	0.297008	0.912045	0.303038	0.287489	0.03607
<b>perimeter3</b>	0.079986	0.782914	0.965137	0.358040	0.970387	0.959120	0.23885
<b>area3</b>	0.107187	0.733825	0.941082	0.343546	0.941550	0.959213	0.20671
<b>smoothness3</b>	0.010338	0.421465	0.119616	0.077503	0.150549	0.123523	0.80532
<b>compactness3</b>	-0.002968	0.590998	0.413463	0.277830	0.455774	0.390410	0.47246
<b>concavity3</b>	0.023203	0.659610	0.526911	0.301025	0.563879	0.512606	0.43492
<b>concave_points3</b>	0.035174	0.793566	0.744214	0.295316	0.771241	0.722017	0.50305
<b>symmetry3</b>	-0.044224	0.416294	0.163953	0.105008	0.189115	0.143570	0.39430
<b>fractal_dimension3</b>	-0.029866	0.323872	0.007066	0.119205	0.051019	0.003738	0.49931


32 rows × 32 columns



In [173]: *### Exploratory Data Analysis - display heatmap matrix*  
*# Looking at the heatmap it also appears the following features are colinear*  
*# due to their high correlation values: radius1 perimeter1, area1, radius3, pe*  
*rimeter3, and area3.*

```
plt.subplots(figsize=(20,15))
sns.heatmap(df.corr(), annot=True, square=True, xticklabels=True, cbar=True)
```

Out[173]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1b0d4b56550>



```

In [174]: ### Models - Strongest Indicator
# Using only concave_points3 the strongest indicator as the only indicator we
can get an R-squared value of 0.72

X = df.drop('Diagnosis', axis=1)
y = df['Diagnosis']

x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
best_r_squared = 0
max_deg = 10
formula = 'Diagnosis ~ concave_points3'
model = smf.ols(formula, data=df).fit()

print("Degree: 1", model.rsquared)

for deg in range(2, max_deg + 1):
    formula += f' + np.power(concave_points3, {deg})'
    model = smf.ols(formula, data=df).fit()
    print("Degree: ", deg, model.rsquared)
    if(model.rsquared_adj > best_r_squared):
        best_formula = formula
        best_model = model
        best_r_squared = best_model.rsquared_adj
        best_degree = deg

formula_strongest_ind = 'Diagnosis ~ concave_points3 + np.power(concave_points
3,2) + np.power(concave_points3, 3)'

```

```

Degree: 1 0.6297470235614584
Degree: 2 0.6309352699330446
Degree: 3 0.7033610378146713
Degree: 4 0.7033900659916137
Degree: 5 0.7202997951788024
Degree: 6 0.7204212156055667
Degree: 7 0.7242914692809307
Degree: 8 0.7242967644841682
Degree: 9 0.725312356224284
Degree: 10 0.7253858093393647

```

```

In [175]: ### Models - Multilinear
# Using backwards elimination of features and additional volume features, we can get an R-squared value of 0.80
# Added volume features that may be impactful.

df['AR1'] = df['radius1']*df['area1']
df['AR2'] = df['radius2']*df['area2']
df['AR3'] = df['radius3']*df['area3']
df['R1cube'] = df['radius1']*df['radius1']*df['radius1']
df['R2cube'] = df['radius2']*df['radius2']*df['radius2']
df['R3cube'] = df['radius3']*df['radius3']*df['radius3']
df['strong_ind'] = np.power(df['concave_points3'], 3)

features = list(df.columns)
features.remove('Diagnosis')
features.remove('ID')

formula = f"Diagnosis ~ {' + '.join(features)}"
psiglevel = 0.01
while True:
    model = smf.ols(formula, data=df).fit()
    max_pvalue = model.pvalues[1:].max()
    if max_pvalue > psiglevel and model.pvalues[1:].idxmax()[0] != 'Intercept':
        feature_to_remove = model.pvalues[1:].idxmax()[0]
        features.remove(feature_to_remove)
        formula = f"Diagnosis ~ {' + '.join(features)}"
    else:
        break

Backwards_Formula = formula
print("Backwards Formula:", Backwards_Formula)
print("Rsquared:", model.summary())
print("Rsquared:", model.rsquared)

```



Backwards Formula: Diagnosis ~ perimeter1 + area1 + concavity1 + smoothness2 + concavity2 + concave\_points2 + radius3 + texture3 + compactness3 + symmetry3 + AR1 + R3cube

Rsquared:

OLS Regression Results

```
=====
=
Dep. Variable:          Diagnosis    R-squared:                0.80
0
Model:                  OLS         Adj. R-squared:           0.79
6
Method:                 Least Squares    F-statistic:              185.
4
Date:                  Tue, 12 Dec 2023    Prob (F-statistic):       2.10e-18
5
Time:                  23:19:56          Log-Likelihood:           64.12
1
No. Observations:      569             AIC:                     -102.
2
Df Residuals:          556             BIC:                     -45.7
7
Df Model:              12
Covariance Type:       nonrobust
=====
```

```
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept          0.2941      0.281        1.046      0.296      -0.258
0.846
perimeter1        -0.0741      0.007       -10.243     0.000      -0.088
-0.060
area1              0.0074      0.001         7.117     0.000       0.005
0.009
concavity1         3.6420      0.394         9.247     0.000       2.868
4.416
smoothness2       10.9796      3.957         2.775     0.006       3.208
18.752
concavity2        -5.6529      0.680        -8.311     0.000      -6.989
-4.317
concave_points2   13.7531      2.874         4.785     0.000       8.108
19.398
radius3            0.1708      0.022         7.861     0.000       0.128
0.213
texture3           0.0092      0.002         5.420     0.000       0.006
0.013
compactness3       0.5010      0.125         3.999     0.000       0.255
0.747
symmetry3          0.7909      0.201         3.937     0.000       0.396
1.185
AR1               -0.0001      2.29e-05       -4.659     0.000      -0.000   -6.
16e-05
R3cube            -0.0001      1.15e-05       -8.862     0.000      -0.000   -7.
94e-05
=====
```

```
=====
=
Omnibus:              55.209    Durbin-Watson:           1.73
```

```
9
Prob(Omnibus):          0.000   Jarque-Bera (JB):          75.01
5
Skew:                   0.730   Prob(JB):           5.14e-1
7
Kurtosis:               4.017   Cond. No.           7.35e+0
6
=====
=
```

**Warnings:**

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 7.35e+06. This might indicate that there are

strong multicollinearity or other numerical problems.

Rsquared: 0.800076491433471

```

In [176]: ### Models - Strongest Indicators
# Using only concave_points3 the strongest indicator as the only indicator we
can get an R-squared value of 0.74

features = list(df.columns)
features.remove('Diagnosis')
features.remove('ID')

best = ['',0]
best_features = []
x_train, x_test = train_test_split(df, train_size=0.2, random_state=10)

for p in features:
    model = smf.ols(formula='Diagnosis~'+p, data=x_train).fit()
    if model.rsquared>best[1]:
        best = [p, model.rsquared]

best_features.append(best[0])
features.remove(best[0])

for i in range(4):

    for p in features:
        formula = f'Diagnosis ~ {" + ".join(best_features)} + {p}'
        model = smf.ols(formula, data=x_train).fit()
        if model.rsquared>best[1]:
            best = [p, model.rsquared]
    best_features.append(best[0])
    features.remove(best[0])

formula_best_features = f'Diagnosis ~ {" + ".join(best_features)} + {p}'
print("Best Features formula: ", formula_best_features )
print("Rsquared:", model.summary())
print("Rsquared:", model.rsquared)

```

Best Features formula: Diagnosis ~ concave\_points3 + texture3 + radius3 + R3cube + smoothness3 + strong\_ind

Rsquared:

OLS Regression Results

```
=====
=
Dep. Variable:          Diagnosis    R-squared:                0.74
9
Model:                  OLS         Adj. R-squared:           0.73
7
Method:                 Least Squares    F-statistic:              63.7
2
Date:                  Tue, 12 Dec 2023    Prob (F-statistic):       1.66e-3
0
Time:                  23:19:57          Log-Likelihood:           -1.170
8
No. Observations:      113             AIC:                      14.3
4
Df Residuals:          107             BIC:                      30.7
1
Df Model:              5
Covariance Type:       nonrobust
=====
```

```
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept        -1.5602      0.217      -7.206      0.000      -1.989
-1.131
concave_points3   3.0656      1.086       2.822      0.006       0.912
5.219
texture3          0.0171      0.004       4.162      0.000       0.009
0.025
radius3           0.0802      0.018       4.419      0.000       0.044
0.116
R3cube           -3.939e-05    1.18e-05     -3.333      0.001     -6.28e-05    -
1.6e-05
strong_ind       10.9234     14.662       0.745      0.458     -18.143
39.990
=====
```

```
=====
=
Omnibus:           3.015    Durbin-Watson:           2.24
2
Prob(Omnibus):     0.221    Jarque-Bera (JB):           2.99
7
Skew:              0.356    Prob(JB):                  0.22
4
Kurtosis:          2.638    Cond. No.                   5.67e+0
6
=====
=
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.67e+06. This might indicate that there are

strong multicollinearity or other numerical problems.  
Rsquared: 0.7485915336935886

```

In [178]: # Results
x_train, x_test = train_test_split(df, train_size=0.2, random_state=10)

train_model_strongest_ind = smf.ols(formula=formula_strongest_ind, data=x_train).fit()
model_strongest_ind = smf.ols(formula=formula_strongest_ind, data=x_test).fit()
y_pred_strongest_ind = train_model_strongest_ind.predict(x_test)
fpr_strongest_ind, tpr_strongest_ind, _ = roc_curve(x_test['Diagnosis'], y_pred_strongest_ind)
roc_auc_strongest_ind = auc(fpr_strongest_ind, tpr_strongest_ind)
y_pred_binary_strongest_ind = (y_pred_strongest_ind > 0.5).astype(int)
f1_strongest_ind = f1_score(x_test['Diagnosis'], y_pred_binary_strongest_ind)

print("Strongest Ind - formula: ", formula_strongest_ind)
print("Strongest Ind - ADJ Rsquared:", model_strongest_ind.rsquared_adj)
print("Strongest Ind - ROC-AUC Strongest Ind:", roc_auc_strongest_ind)
print("Strongest Ind - F1 Score Strongest Ind:", f1_strongest_ind)

train_model_backwards = smf.ols(formula=Backwards_Formula, data=x_train).fit()
model_backwards = smf.ols(formula=Backwards_Formula, data=x_test).fit()
y_pred_backwards = train_model_backwards.predict(x_test)
fpr_backwards, tpr_backwards, _ = roc_curve(x_test['Diagnosis'], y_pred_backwards)
roc_auc_backwards = auc(fpr_backwards, tpr_backwards)
y_pred_binary_backwards = (y_pred_backwards > 0.5).astype(int)
f1_backwards = f1_score(x_test['Diagnosis'], y_pred_binary_backwards)

print('\n')
print("Backwards Refinement - formula: ", Backwards_Formula)
print("Backwards Refinement - ADJ Rsquared:", model_backwards.rsquared_adj)
print("Backwards Refinement - ROC-AUC Backwards:", roc_auc_backwards)
print("Backwards Refinement - F1 Score Backwards:", f1_backwards)

train_model_best_features = smf.ols(formula=formula_best_features, data=x_train).fit()
model_best_features = smf.ols(formula=formula_best_features, data=x_test).fit()
y_pred_best_features = train_model_best_features.predict(x_test)
fpr_best_features, tpr_best_features, _ = roc_curve(x_test['Diagnosis'], y_pred_best_features)
roc_auc_best_features = auc(fpr_best_features, tpr_best_features)
y_pred_binary_best_features = (y_pred_best_features > 0.5).astype(int)
f1_best_features = f1_score(x_test['Diagnosis'], y_pred_binary_best_features)

print('\n')
print("Forward Refinement - formula: ", formula_best_features)

```

```
print("Forward Refinement - ADJ Rsquared:", model_best_features.rsquared_adj)
print("Forward Refinement - ROC-AUC Best Features:", roc_auc_best_features)
print("Forward Refinement - F1 Score Best Features:", f1_best_features)
```

Strongest Ind - formula: Diagnosis ~ concave\_points3 + np.power(concave\_points3,2) + np.power(concave\_points3, 3)  
 Strongest Ind - ADJ Rsquared: 0.6978324367912714  
 Strongest Ind - ROC-AUC Strongest Ind: 0.960978835978836  
 Strongest Ind - F1 Score Strongest Ind: 0.8711656441717791

Backwards Refinement - formula: Diagnosis ~ perimeter1 + area1 + concavity1 + smoothness2 + concavity2 + concave\_points2 + radius3 + texture3 + compactness3 + symmetry3 + AR1 + R3cube  
 Backwards Refinement - ADJ Rsquared: 0.7915645212265836  
 Backwards Refinement - ROC-AUC Backwards: 0.9954117063492064  
 Backwards Refinement - F1 Score Backwards: 0.9473684210526314

Forward Refinement - formula: Diagnosis ~ concave\_points3 + texture3 + radius3 + R3cube + smoothness3 + strong\_ind  
 Forward Refinement - ADJ Rsquared: 0.7402375060491297  
 Forward Refinement - ROC-AUC Best Features: 0.9908027447089948  
 Forward Refinement - F1 Score Best Features: 0.9415384615384617

In [ ]: *## Discussion and Conclusion*

```
# Best r_squared value value is backwards refinement. It looks like even just
using the polynomial concatenation of the
# strongest indicator results in a 0.96 ROC AUC

# Best formula was using the backwards refinement which resulted in an adjusted
R2 of 0.79 and an ROC-AUC curve of 0.995

# Some key take-aways is that with even just the strongest indicator you can get
some accurate classification of a cancerous tumor.
# In the future, this can generate an even more accurate model using a neural
network.
```

In [ ]:

In [ ]: